

# Query Optimization Costs-II

# Knowledge objectives

1. Enumerate five join algorithms
2. Explain the prerequisites of each join algorithm (i.e. Row Nested Loops, Block Nested Loops, Hash Join, Sort-Match)
3. Write the pseudo-code of five join algorithms

# Understanding objectives

1. Identify when a join algorithm can or cannot be used in a given query or process tree
2. Given the statistics of the tables, estimate the cost of a join using each of the algorithms

# Clustered structure

# Clustered Structure

- Space
  - $\lceil 1.5B_{RS} \rceil$
- Access paths
  - Cost of table scan of R
    - $\lceil 1.5B_{RS} \rceil$
  - Cost of table scan of S
    - $\lceil 1.5B_{RS} \rceil$

$$B_{RS} = B_R + B_S$$
$$R_{RS} = (|R| + |S|) / (B_R + B_S)$$

Dept 1	Employee 14	Employee 8	Employee 6		Dept 2	Employee 3			Dept 3		Dept 4	Employee 18	Employee 2	
--------	-------------	------------	------------	--	--------	------------	--	--	--------	--	--------	-------------	------------	--

# Nested Loops

# Row Nested Loops

- Algorithm
  - for each block of R
    - read block of R into a memory page
    - for each tuple t in the read page
      - if there is an index in S for attribute A
        - go through the index of S.A using the value of t.A
        - if there is any tuple satisfying the join condition
          - if we are interested in attributes of S
            - go to the corresponding tuples of S
          - endiff
          - generate result
        - endiff
      - else scan the whole table S and generate result
    - endiff
  - endForEach
  - endForEach
- Cost, if we do NOT look for attributes of S (semi-join)
  - B+:  $B_R + |R| \cdot (h_S + (k-1)/u_S)$
  - Hash:  $B_R + |R|$
- Cost, if we DO look for attributes of S
  - B+:  $B_R + |R| \cdot (h_S + (k-1)/u_S + k)$
  - Clustered:  $B_R + |R| \cdot (h_S + 1 + (k-1)/((2/3)R_S))$
  - Hash:  $B_R + |R| \cdot (1 + k)$
- Considerations
  - It is only useful if there is an index over the join attribute
    - The table with the index is the internal one
      - If both have an index, we can choose which is the internal one (we must compare both costs)
      - If, according to the process tree, we already performed some operation over the table, its index cannot be used anymore
  - A hash index can only be used for equi-join
  - $B_R$  is the real number of blocks (taking into account a possible cluster)

$$u = \%load \cdot 2d = (2/3) \cdot 2d$$
$$h = \lceil \log_u |T| \rceil - 1$$

k = average appearance of each value of R.A in S.A

# Block Nested Loops (I)

- Algorithm

repeat

    read M blocks of R

for each block of S

        read block of S into a memory page

for each tuple t in the pages of R

for each tuple s in the page of S

if (t.A  $\theta$  s.B) then generate result

endif

endForEach

endForEach

endForEach

until no more blocks to read from R

- Cost (with  $M+2$  memory pages)

- $B_R + B_S \cdot \lceil B_R/M \rceil$

- Considerations

- It can always be used (even for  $\theta$ -join)

- It is of special interest if  $B_R \leq M$

- It is not symmetric

- It is better if the smaller table is in the external loop

- $B_R$  and  $B_S$  are the real number of blocks (taking into account a possible cluster)

# Block Nested Loops (II)

M=2

E	E	I	O
R <sub>3</sub>	R <sub>4</sub>	S <sub>1</sub>	

R<sub>1</sub> R<sub>2</sub> R<sub>3</sub> R<sub>4</sub>

S<sub>1</sub> S<sub>2</sub>



Accesses

R R R R R R R R R R

E	E	I	O
R <sub>1</sub>	R <sub>2</sub>	S <sub>4</sub>	

R<sub>1</sub> R<sub>2</sub>

S<sub>1</sub> S<sub>2</sub> S<sub>3</sub> S<sub>4</sub>

# Hash-Join

One-pass

Two-pass



# One-pass hash Join

- Algorithm
  - Build phase
    - Create an empty hash table
    - for each block of S
      - read block of S into a memory page
      - for each tuple t in the read page
        - register it in the hash table under key  $h(t.A)$
      - endForEach
    - endForEach
  - Probe phase
    - for each block of R
      - read block of R into a memory page
      - for each tuple t in the read page
        - go to the hash table using the key  $h(t.A)$
        - if there is any tuple satisfying the join condition
          - generate result
        - endif
      - endForEach
    - endForEach
- Cost, with  $M+2$  memory pages
  - $B_R + B_S$
- Considerations
  - A hash index is not required to exist a priori
    - It is built in the first phase
  - It can only be used for equi-join
  - It is not symmetric
    - It chooses the smaller table to build the hash table (hence, it is in the internal loop during the probing phase)
  - It requires  $B_{\text{Smaller}} \leq M$  (so that its hash index fits in memory)

# Two-pass hash Join (I)

- Algorithm

- Partition R into  $p$  parts using a hash function
  - Partition S into  $p$  parts using the same hash function
  - Use One-pass hash join  $p$  times (part by part)

- Cost, with  $M+2$  memory pages

- $2B_R + 2B_S + B_R + B_S$ 
    - $B_R$  and  $B_S$  are the real number of blocks (taking into account a possible cluster)
      - Beware that after partitioning, the tables have no empty space anymore

- Considerations

- We must guarantee that each part of S fits into M memory pages later, so we take  $p = \lceil B_{\text{Smaller}}/M \rceil$ 
    - Thus, the size of each part  $B_{\text{Smaller}}/p$  is M pages
      - Assuming the hash function results in a uniform distribution of values
  - Assumes  $M < B_{\text{Smaller}} \leq M^2 + M$ 
    - In this case  $1 < p \leq M + 1$  (notice that  $\lceil M/M \rceil = 1$  and  $\lceil (M^2 + M)/M \rceil = M + 1$ )
      - If  $B_{\text{Smaller}} \leq M$  (i.e.,  $p = 1$ ), then we use One-pass
      - If  $B_{\text{Smaller}} > M^2 + M$  (i.e.,  $p > M + 1$ ), then Sort-Match will be preferred

# Two-pass hash Join (II)

Accesses

R
R
W
W
W
R
W
R
W
<u>Pass 1</u>
R
R
R
R
<u>Pass 2</u>

M=1

B<sub>1</sub>=B<sub>2</sub>=2

p=2




Table 1

O	E
1	2
3	2
5	3

Table 2 (hash)

O	E
5	2
7	4
2	4

# Sort-Match

# Sort-Match (I)

- Algorithm (for equi-join on an attribute A which is unique, at least, in R)

```
Sort R and S (if necessary)
tR:=first(R); tS:=first(S);
while not (end(R) or end(S))
    if (tR[A]<tS[A]) tR:=next(R);
    elsif (tR[A]>tS[A]) tS:=next(S);
    else generate result from tR and tS; tS:=next(S);
    endif
endWhile
```

- Cost, with M+1 memory pages

- Sorting R (same for S):
  - If sorted: 0
  - Elsif  $B_R \leq M$ :  $2B_R$
  - Else:  $2B_R \lceil \log_M B_R \rceil$
- Merging R and S:  $B_R + B_S$


- Considerations

- It is only useful for equi-join, <-join, >-join and anti-join
  - The given cost corresponds to equi-join and anti-join
- It is of special interest when at least one table is Clustered (or somehow sorted by the join attribute)
  - If both tables are pre-sorted, we only need 3 memory pages
- The result is already sorted
- $B_R$  and  $B_S$  are the real number of blocks (taking into account a possible cluster)
  - Beware that after sorting, the auxiliary tables have no empty space anymore

# Sort-Match (II)

Accesses

R  
R  
R  
R




# Example of cost-based optimization

# Cost-based optimization steps

1. Generate alternatives in the search space
  - a. Join order
  - b. Potential algorithms
  - c. Available structures (access path)
  - d. ~~Materialization or not of intermediate results~~
    - We will assume that they are always materialized
2. Evaluate those alternatives
  - a. Intermediate results cardinality and size estimation
  - b. Cost estimation
3. Choose the best option
4. Generate the corresponding access plan

# Example of cost-based optimization (I)


- The tables have the following structures:
  - Producers
    - Clustered by prodId
    - $B^+$  by region
  - Wines
    - Clustered by winId
  - Vintages
    - Clustered by winId and prodId
- We have the following statistics:
  - Tables (extra space due to being clustered needs to be added)
    - $|P|=10000 R_p=12 B_p=834$
    - $|W|= 5000 R_w=10 B_w=500$
    - $|V|= 100000 R_v=20 B_v=5000$
  - Attributes
    - prodId, winId and strength: length=5 bytes
    - ndist(region)=30
    - min(quantity)=10 max(quantity)=500
    - ndist(strength)=100
- Moreover, we know that
  - There are 500 useful bytes per intermediate disk block
  - Each table is in a different file (there is no Clustered Structure)
  - Cost of accessing disk blocks is 1 second ( $D=1$ )
  - Cost of CPU processing is negligible ( $C=0$ )
  - The order of B-trees is 75
  - The DBMS can use:
    - Block Nested Loops (with 6 memory pages,  $M=4$ )
    - Row Nested Loops
    - Sort Match (with 3 memory pages for sorting,  $M=2$ )
- We will not change the order of operations coming from syntactic optimization



# Example of cost-based optimization (II)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- Step 2b: Cost estimation for each algorithm
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength  
FROM wines w, producers p, vintages v  
WHERE v.winId=w.winId  
    AND p.prodId=v.prodId  
    AND p.region="Priorat"  
    AND v.quantity>100;
```



# Example of cost-based optimization (III)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- Step 2b: Cost estimation for each algorithm
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength
FROM wines w, producers p, vintages v
WHERE v.winId=w.winId
      AND p.prodId=v.prodId
      AND p.region="Priorat"
      AND v.quantity>100;
```

## PT1/PT2

### ▪ Selection over V: V'

Record length  $V' = 5+5=10$  bytes

$SF(quantity > 100) =$

$$= (\max(quantity) - 100) / (\max(quantity) - \min(quantity)) =$$

$$= 0.81632$$

$$|V'| = SF * |V| = 0.81632 * 100,000 = 81,632$$

$$R_V = \lfloor 500/10 \rfloor = 50 \text{ records/block}$$

$$B_V = \lceil 81,632/50 \rceil = 1,633 \text{ blocks}$$

### ▪ Selection over P: P'



Record length  $P' = 5$  bytes

$SF(region = "Priorat") = 1/n_{dist}(region) = 1/30$

$$|P'| = SF * |P| = 10000/30 = 333$$

$$R_P = \lfloor 500/5 \rfloor = 100 \text{ records/block}$$

$$B_P = \lceil 333/100 \rceil = 4 \text{ blocks}$$



# Example of cost-based optimization (IV)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- Step 2b: Cost estimation for each algorithm
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength  
FROM wines w, producers p, vintages v  
WHERE v.winId=w.winId  
AND p.prodId=v.prodId  
AND p.region="Priorat"  
AND v.quantity>100;
```

## PT1

### ■ Join between W and V': WV'


Record length WV' = 5+5 bytes

$$SF = 1/|W| = 1/5000$$

$$|WV'| = SF * |W| * |V'| = |V'| = 81,632$$

$$R_{WV'} = \lfloor 500/10 \rfloor = 50 \text{ records/block}$$

$$B_{WV'} = \lceil 81,632/50 \rceil = 1,633 \text{ blocks}$$



### ■ Join between WV' and P': WV'P' (if quantity and region are independent)


Record length WV'P' = 5 bytes

$$SF(WV'*P') = (1/|P'|)*(1/30) = 10^{-4}$$

$$|WV'P'| = SF * |WV'| * |P'| = 10^{-4} * |WV'| * |P'| = 2,721$$

$$R_{WV'P'} = \lfloor 500/5 \rfloor = 100 \text{ records/block}$$

$$B_{WV'P'} = \lceil 2721/100 \rceil = 28 \text{ blocks}$$



# Example of cost-based optimization (V)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- Step 2b: Cost estimation for each algorithm
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength  
FROM wines w, producers p, vintages v  
WHERE v.winId=w.winId  
      AND p.prodId=v.prodId  
      AND p.region="Priorat"  
      AND v.quantity>100;
```

## PT2

### ■ Join between V' and P': V'P' (if quantity and region independent)


Record length V'P' = 5 bytes

$$SF(V' * P') = (1/30) * (1/|P'|) = 10^{-4}$$

$$|V'P'| = SF * |V'| * |P'| = 10^{-4} * |V'| * |P'| = 2,721$$

$$R_{V'P'} = \lfloor 500/5 \rfloor = 100 \text{ records/block}$$

$$B_{V'P'} = \lceil 2721/100 \rceil = 28 \text{ blocks}$$



### ■ Join between W and V'P': WV'P'


Record length WV'P' = 5 bytes

$$SF = 1/|W|$$

$$|WV'P'| = SF * |W| * |V'P'| = |V'P'| = 2,721$$

$$R_{WV'P'} = \lfloor 500/5 \rfloor = 100 \text{ records/block}$$

$$B_{WV'P'} = \lceil 2721/100 \rceil = 28 \text{ blocks}$$



# Example of cost-based optimization (VI)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- Step 2b: Cost estimation for each algorithm
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength  
FROM wines w, producers p, vintages v  
WHERE v.winId=w.winId  
      AND p.prodId=v.prodId  
      AND p.region="Priorat"  
      AND v.quantity>100;
```

## PT1/PT2

### ■ Final result: O

$$\begin{aligned} \text{Record length } O &= 5 \text{ bytes} \\ |O| &= \text{ndist(strength)} = 100 \\ R_o &= \lfloor 500/5 \rfloor = 100 \text{ records/block} \\ B_o &= \lceil 100/100 \rceil = 1 \text{ blocks} \end{aligned}$$

# Example of cost-based optimization (VII)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- **Step 2b: Cost estimation for each algorithm**
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength
FROM wines w, producers p, vintages v
WHERE v.winId=w.winId
      AND p.prodId=v.prodId
      AND p.region="Priorat"
      AND v.quantity>100;
```

## PT1/PT2

### ■ Selection over V: V'

Available access paths: No index

$$\text{cost}_{\text{Scan}}(V') = \lceil 1.5B_V \rceil = \lceil 1.5 * 5,000 \rceil = 7,500$$

**Choose Scan**

### ■ Selection over P: P'

Available access paths: B+ and No index

$$\text{cost}_{\text{Scan}}(P') = \lceil 1.5 * B_P \rceil = \lceil 1.5 * 834 \rceil = 1,251$$

$$\text{cost}_{B+}(P') = \lceil \log_{100} |P| \rceil - 1 + \text{SF}(\text{region}=\text{"Priorat"}) * |P| + ((\text{SF}(\text{region}=\text{"Priorat"})) * |P| - 1) / 100 = 1 + 333 + 332 / 100 = 337$$

**Choose B+**

### ■ Sort of WV'P': O

$$\text{cost}_{\text{MergeSort}}(O) = 2B_{WV'P'} \cdot \lceil \log_M(B_{WV'P'}) \rceil - B_{WV'P'} = 2 \cdot 28 \cdot \lceil \log_2(28) \rceil - 28 = 252$$

# Example of cost-based optimization (VIII)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- **Step 2b: Cost estimation for each algorithm**
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength  
FROM wines w, producers p, vintages v  
WHERE v.winId=w.winId  
      AND p.prodId=v.prodId  
      AND p.region="Priorat"  
      AND v.quantity>100;
```

## PT1

### ■ Join between W and V': WV'

Available algorithms:

Block Nested Loops

$\lceil 1.5 \cdot B_W \rceil < B_{V'}$  (use commutative property of joins)

$$\text{cost}_{\text{NestedLoop}}(\text{WV}') = \lceil 1.5B_W \rceil + \lceil 1.5B_W/M \rceil * B_{V'} = \lceil 1.5*500 \rceil + \lceil 1.5*500/4 \rceil * 1633 = 307,754$$

Row Nested Loops

Yes, we do look for attributes of W

V' does not use extra space any more for being ordered

$$\text{cost}_{\text{RowNestedLoops}}(\text{WV}') = B_{V'} + |V'| * (\lceil \log_{100} |W| \rceil - 1 + 1 + (1.5(k-1)/10)) = 1,633 + 81,632 * (\lceil \log_{100} 5,000 \rceil - 1 + 1) = 164,897$$

Sort-Match

W is ordered by winID, V' is still ordered by winId and prodId

$$\text{cost}_{\text{SortMatch}}(\text{WV}') = \lceil 1.5B_W \rceil + B_{V'} = \lceil 1.5*500 \rceil + 1,633 = 2,383$$

Choose Sort-Match

# Example of cost-based optimization (IX)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- **Step 2b: Cost estimation for each algorithm**
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength
FROM wines w, producers p, vintages v
WHERE v.winId=w.winId
      AND p.prodId=v.prodId
      AND p.region="Priorat"
      AND v.quantity>100;
```

## PT1

### ■ Join between WV' and P': WV'P'

Available algorithms:

Block Nested Loops

$B_{P'} < B_{WV'}$  (use commutative property of joins)

$$\text{cost}_{\text{NestedLoop}}(\text{WV}'\text{P}') = B_{P'} + \lceil B_{P'}/M \rceil * B_{WV'} = 4 + \lceil 4/4 \rceil * 1,633 = 1,637$$

Sort Match

Neither WV' nor P' are ordered by prodId

$$\text{cost}_{\text{SortMatch}}(\text{WV}'\text{P}') = 2*B_{WV'}*\lceil \log_2 B_{WV'} \rceil + 2*B_{P'}*\lceil \log_2 B_{P'} \rceil + B_{WV'} + B_{P'} = 2*1,633*11 + 2*4*2 + 1633 + 4 = 37,579$$

Choose Nested Loops

# Example of cost-based optimization (X)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- **Step 2b: Cost estimation for each algorithm**
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength
FROM wines w, producers p, vintages v
WHERE v.wineId=w.wineId
      AND p.productId=v.productId
      AND p.region="Priorat"
      AND v.quantity>100;
```

## PT2

### ■ Join between V' y P': V'P'

Available algorithms:

Block Nested Loops

$B_{P'} < B_{V'}$  (use commutative property of joins)

$$\text{cost}_{\text{NestedLoop}}(V'P') = B_{P'} + \lceil B_{P'} / M \rceil * B_{V'} = 4 + \lceil 4/4 \rceil * 1,633 = 1,637$$

Sort Match

Neither V' nor P' are ordered by prodId

$$\text{cost}_{\text{SortMatch}}(V'P') = 2 * B_{V'} * \lceil \log_2 B_{V'} \rceil + 2 * B_{P'} * \lceil \log_2 B_{P'} \rceil + B_{V'} + B_{P'} = 2 * 1,633 * 11 + 2 * 4 * 2 + 1,633 + 4 = 37,579$$

**Choose Nested Loops**

# Example of cost-based optimization (XI)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- **Step 2b: Cost estimation for each algorithm**
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength
FROM wines w, producers p, vintages v
WHERE v.winId=w.winId
      AND p.prodId=v.prodId
      AND p.region="Priorat"
      AND v.quantity>100;
```

## PT2

### ■ Join between W y V'P': WV'P'

Available algorithms:

Block Nested Loops

$B_{V'P'} < \lceil 1.5B_W \rceil$  (use commutative property of joins)

$$\text{cost}_{\text{NestedLoop}}(\text{WV}'\text{P}') = B_{V'P'} + \lceil B_{V'P'} / M \rceil * \lceil 1.5B_W \rceil = 28 + \lceil 28/4 \rceil * \lceil 1.5 * 500 \rceil = 5278$$

Row Nested Loops

Yes, we look for attributes of W

$$\text{cost}_{\text{RowNestedLoops}}(\text{WV}'\text{P}') = B_{V'P'} + |V'P'| * (\lceil \log_{100} |W| \rceil - 1 + 1 + (1.5(k-1)/10)) = 28 + 2,721 * (\lceil \log_{100} 5,000 \rceil - 1 + 1) = 5,470$$

Sort-Match

W is sorted by winId, V'P' is not sorted by winId


$$\text{cost}_{\text{SortMatch}}(\text{WV}'\text{P}') = 2B_{V'P'} \lceil \log_2 B_{V'P'} \rceil + \lceil 1.5B_W \rceil + B_{V'P'} = 2 * 28 * \lceil \log_2 28 \rceil + \lceil 1.5 * 500 \rceil + 28 = 1,058$$

**Choose Sort-Match**

# Example of cost-based optimization (XII)

- Step 1: Generate alternatives
- Step 2a: Intermediate results estimation
- Step 2b: Cost estimation for each algorithm
- Step 3: Choose the best option

```
SELECT DISTINCT w.strength
FROM wines w, producers p, vintages v
WHERE v.winId=w.winId
AND p.prodId=v.prodId
AND p.region="Priorat"
AND v.quantity>100;
```



# Closing

# Summary table

	No index	B+	Hash	Clustered	Clustered structure
All tuples	Scan				
One tuple		Go through index Go to table	Apply function Go to bucket Go to table	Go through index Go to table	
Several tuples		Go through index Follow leaves Go to table		Go through index Go to table Scan table	
Join	Block Nested Loops Or Hash Join	Row Nested Loops	Row Nested Loops	Row Nested Loops Or Sort-Match	Scan

# Summary

- Cost estimation
  - Join algorithms
    - Clustered structure
    - Nested loops
      - Row
      - Block
    - Hash
      - One-pass
      - Two-pass
    - Sort-Match

# Bibliography

- Y. Ioannidis. *Query Optimization*. ACM Computing Surveys, vol. 28, num. 1, March 1996
- G. Gardarin and P. Valduriez. *Relational Databases and Knowledge Bases*. Addison-Wesley, 1998
- J. Sistac. *Sistemes de Gestió de Bases de Dades*. Editorial UOC, 2002
- R. Ramakrishnan and J. Gehrke. *Database Management Systems*. McGraw-Hill, 3<sup>rd</sup> Edition, 2003
- J. Lewis. *Cost-Based Oracle Fundamentals*. Apress, 2006
- S. Lightstone, T. Teorey and T. Nadeau. *Physical Database Design*. Morgan Kaufmann, 2007